

# INSURANCE CLAIMS FRAUD DETECTION

## What are we going to discuss today?

Today we have been offered insurance for almost everything, you know what even for life, e.g. Health Insurance, Property Insurance, Automobile Insurance, even Insurance for mobiles, Life Insurance, Electronic gadgets what-not, the list goes on..., Let us know the challenges faced by the insurance companies.

What even insurance companies have challenges? Why not every business model has its own loop holes and challenges. Let us see the challenges and how can we help them with our knowledge of Machine Learning and Data Science.

## Business Problem:

The most common and biggest problem faced by Insurance Companies is FRAUD CLAIMS, which is a huge loss to company. So we will try to find out whether the claim is Genuine or Fraud based on the data from an Auto Insurance Company by building a Machine Learning model.



Image source: Google – Credits to respected owner created

Before getting started with the project in building Machine Learning model, even before viewing the data it is always best practice to know about the problem / business domain and have some idea of it. Which is helpful in our entire project phase. So I recommend to have a quick check with our best buddy, you're right google the best source for basic insights.

## Glimpse of Automobile Insurance Domain:

In an automobile insurance the company have some certain terms and conditions, checklists where the insurance will be covered and where it doesn't take responsibility or pay any amount if it doesn't meet their terms like : Not eligible for claim in case of accident when drunk and driven, etc.

*The Fraud Claim:* It is a false claim by user even if they don't meet the terms of company but will show as the Genuine and make claim to the company that they followed all their terms and is eligible for claim for their own benefits / to not occur loss.

## DATA SOURCE:

The source of data to me is from DATA TRAINED Institute.

File format: .csv

The data is collection of different parameters and checks with the claim status whether fraud or not collected from different insurance companies and compiled into one file.

## Step by Step process:

### 1. Import the dataset:

Importing the dataset in jupyter notebook and saving to a variable

```
In [1]: #import necessary modules
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: #importing the dataset
data=pd.read_csv('8.Insurance.csv')

#Setting to display all columns
pd.set_option('display.max_columns', None)
```

Cell [1]: we imported some required modules to get data and create visualizations.

8. Insurance.csv is the file name which I stored in my local machine at current working directory.

**Data Analysis:** Narrow method of checking assumptions for model fitting, hypothesis testing by cleaning, filling missing values and making transformations as required.

```
In [3]: #see the data
data.head()
```

Out[3]:

Out[3]:

Out[3]:

We are now seeing the top 5 rows of our dataset. The data we have consists of columns with values categorical and numerical type also few columns with '?' in them.

Let us see the shape and size of dataset:

```
In [4]: data.shape
```

```
Out[4]: (1000, 40)
```

```
In [5]: data.size
```

```
Out[5]: 40000
```

There are 1000 rows and 40 columns, with a size of 40000 values (product of rows and columns)

### Column details: Important step in EDA

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 40 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   months_as_customer                    1000 non-null   int64  
 1   age                                    1000 non-null   int64  
 2   policy_number                         1000 non-null   int64  
 3   policy_bind_date                      1000 non-null   object  
 4   policy_state                          1000 non-null   object  
 5   policy_csl                            1000 non-null   object  
 6   policy_deductable                     1000 non-null   int64  
 7   policy_annual_premium                 1000 non-null   float64 
 8   umbrella_limit                        1000 non-null   int64  
 9   insured_zip                           1000 non-null   int64  
10  insured_sex                           1000 non-null   object  
11  insured_education_level                1000 non-null   object  
12  insured_occupation                     1000 non-null   object  
13  insured_hobbies                        1000 non-null   object  
14  insured_relationship                   1000 non-null   object  
15  capital-gains                          1000 non-null   int64  
16  capital-loss                           1000 non-null   int64  
17  incident_date                          1000 non-null   object  
18  incident_type                          1000 non-null   object  
19  collision_type                         1000 non-null   object  
20  incident_severity                      1000 non-null   object  
21  authorities_contacted                  1000 non-null   object  
22  incident_state                         1000 non-null   object  
23  incident_city                          1000 non-null   object  
24  incident_location                      1000 non-null   object  
25  incident_hour_of_the_day                1000 non-null   int64  
26  number_of_vehicles_involved             1000 non-null   int64  
27  property_damage                        1000 non-null   object  
28  bodily_injuries                        1000 non-null   int64  
29  witnesses                              1000 non-null   int64  
30  police_report_available                 1000 non-null   object  
31  total_claim_amount                     1000 non-null   int64  
32  injury_claim                           1000 non-null   int64  
33  property_claim                         1000 non-null   int64  
34  vehicle_claim                          1000 non-null   int64  
35  auto_make                              1000 non-null   object  
36  auto_model                             1000 non-null   object  
37  auto_year                              1000 non-null   int64  
38  fraud_reported                         1000 non-null   object  
39  _c39                                   0 non-null      float64 
dtypes: float64(2), int64(17), object(21)
memory usage: 312.6+ KB
```

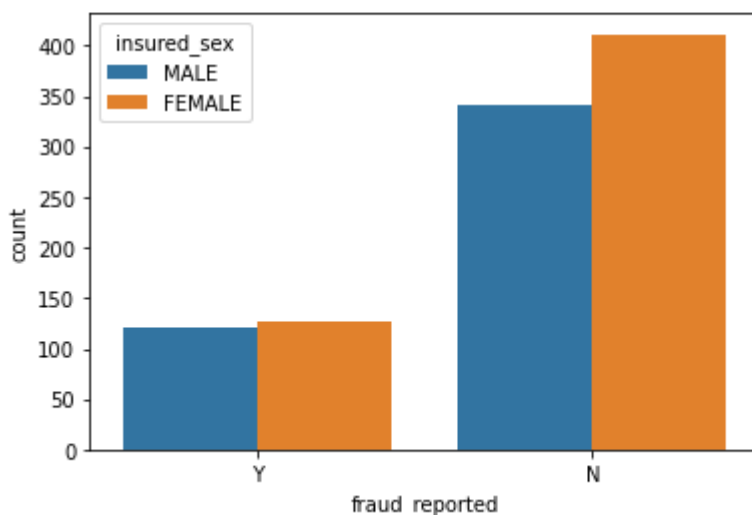
We got the brief of the data columns like – column name, number of non-null values and its data type. This is important step to see whether there are any datatype mismatches and to fix it. So we need to carefully observe the type of data present and verify here.

Observations:

- Our data is mixture of integer, float and object data types.
- Our target variable 'fraud\_reported' is object type. So we use 'Classification Model'.
- We see there are no null values in our brief data, but we found '?' in some columns, so we need to replace them in the best possible way.
- Column policy\_csl is shown as object but we can see it contains numerical data. There might be chances of '?' or any other special character or text present in that column. So it is shown as object type. We need to find that and convert its type.
- Observe column 39: '\_c39' - it has 0 non null values, i.e. all the values are null. So we can drop that column straight away.
- **Assumptions:**
  - Assuming that, for making a fraud it is personal intention and that doesn't depend on what kind of auto make, model and year. So I opt to drop these columns also.
  - Policy\_number: It is a unique id given to the customer, to track the subscription status and other details of customer. So we can drop this column, before dropping check if there are any duplicates in that column, if there are duplicates remove those rows by keeping one uniquely then drop the column.
  - Insured\_zip: It is the zip code where the insurance was made. It doesn't give any information or useful to us. so we will drop this column too.
  - Policy\_csl is basically [Combined Single Limit](#), i.e. how much credible and what their premium covers, it is also not necessary. So we will drop it.

We have seen that fraud reported by male & female both are almost same count.

```
: sns.countplot(data['fraud_reported'], hue=data['insured_sex'])  
C:\Users\Lucky Girish\anaconda3\lib\site-packages\seaborn\_decorators.py:100: FutureWarning:   
keyword arg: x. From version 0.12, the only valid positional argument is 'data'. Using a  
implicit keyword will result in an error or misinterpretation.  
warnings.warn(  
: <matplotlib.axes._subplots.AxesSubplot at 0x22b55456a90>
```



```
In [10]: #Dropping the columns
dfc=dfc.drop(['policy_number','insured_zip','auto_make','auto_model','auto_year','_c39','policy_csl'],axis=1)
```

We dropped the columns and copy of it saved into new variable 'dfc', the initial variable data will have all columns. Where as in dfc we have columns that aren't dropped.

```
In [11]: dfc.columns
```

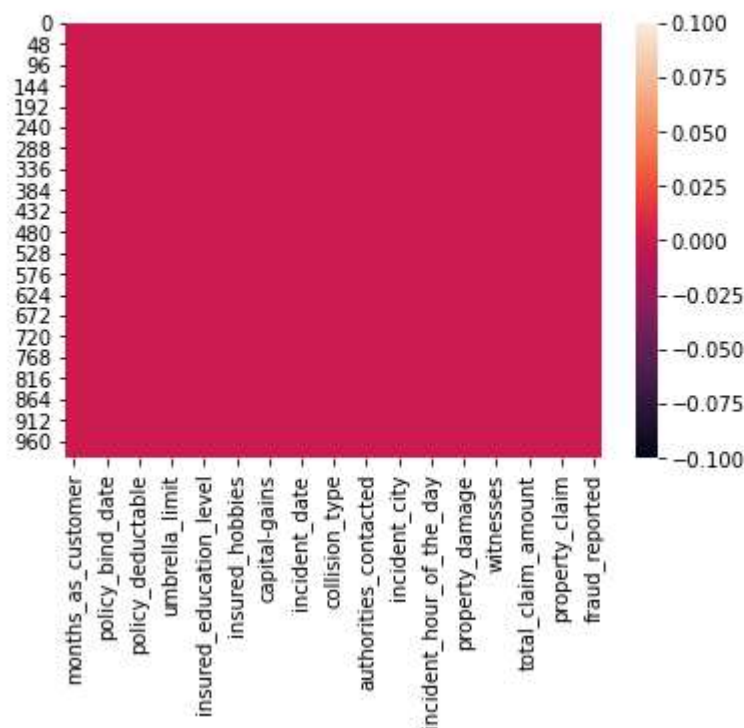
```
Out[11]: Index(['months_as_customer', 'age', 'policy_bind_date', 'policy_state',
               'policy_deductable', 'policy_annual_premium', 'umbrella_limit',
               'insured_sex', 'insured_education_level', 'insured_occupation',
               'insured_hobbies', 'insured_relationship', 'capital-gains',
               'capital-loss', 'incident_date', 'incident_type', 'collision_type',
               'incident_severity', 'authorities_contacted', 'incident_state',
               'incident_city', 'incident_location', 'incident_hour_of_the_day',
               'number_of_vehicles_involved', 'property_damage', 'bodily_injuries',
               'witnesses', 'police_report_available', 'total_claim_amount',
               'injury_claim', 'property_claim', 'vehicle_claim', 'fraud_reported'],
              dtype='object')
```

This gives list of columns present in the variable dfc.

Checking for null values:

```
In [12]: #Visualising null values
sns.heatmap(dfc.isnull())
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x23c74090850>
```



There are no null values. We are going to replace the '?' with mean, median or mode or any other based on the column.

We see '?' present in many columns:

for e.g. we deal with one column collision type-

We can see there are 4 unique values in this column, as shown below. Front, Rear, Side and '?' Collisions. '?' is there in 178 rows. This might not be filled, it is a manual error. Because the collision should be in one of the category, and must note while claimed.

The best, thumb hand rule to deal this kind in case of categorical values is to fill with the mode of column.

So we will replace these '?' with the most repeated type of collision type: Rear Collision

```
In [13]: #Collision_type  
dfc['collision_type'].nunique()
```

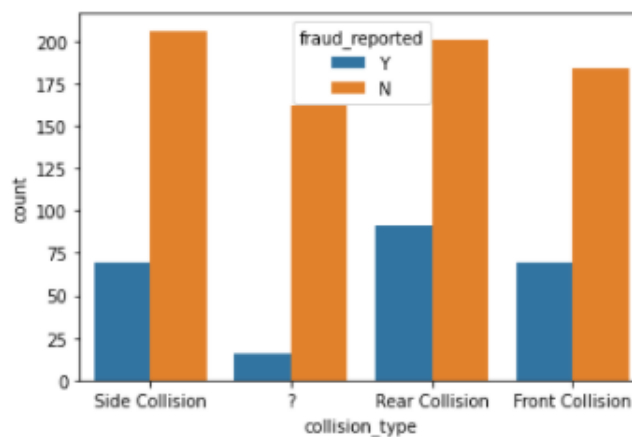
```
Out[13]: 4
```

```
In [15]: dfc['collision_type'].value_counts()
```

```
Out[15]: Rear Collision    292  
Side Collision    276  
Front Collision    254  
?    178  
Name: collision_type, dtype: int64
```

```
In [14]: sns.countplot(dfc['collision_type'],hue=dfc['fraud_reported'])
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x23c74805a30>
```

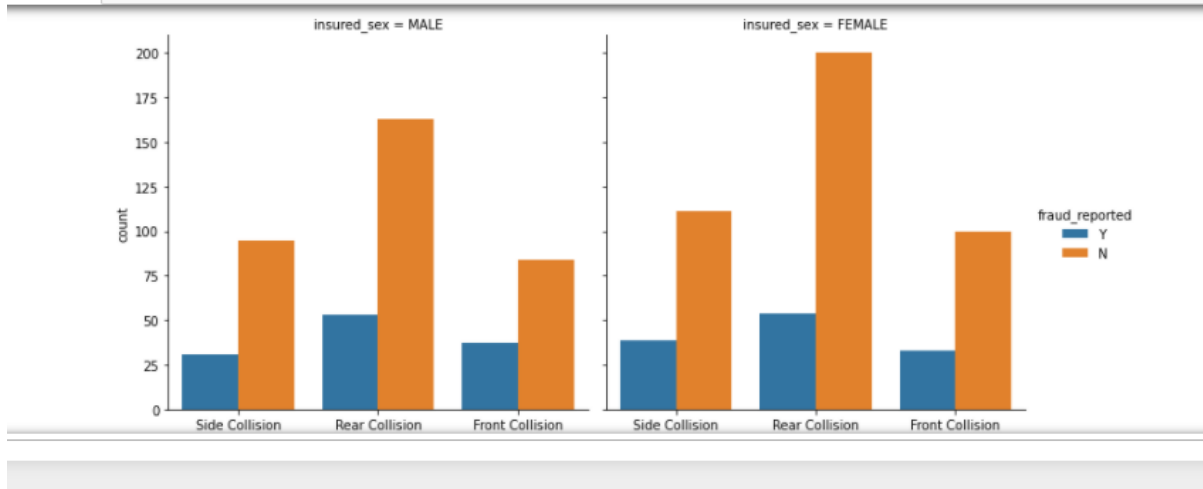




Replacing and counts after replaced wrt to fraud claim or not:

```
In [27]: dfc['collision_type']=dfc['collision_type'].replace('?', 'Rear Collision')
```

```
In [37]: sns.catplot(x='collision_type', hue='fraud_reported', col='insured_sex', kind="count", data=dfc)
```

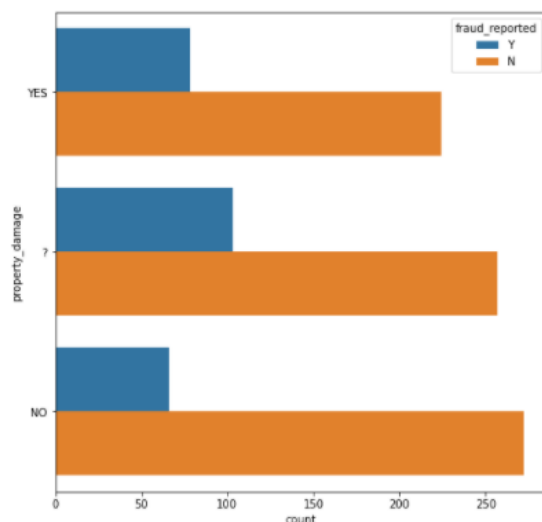


More collisions done by females than males. Rear collisions are max. From above we can also see, what kind of collision made by male or female and whether it is fraud claim or not. Almost count of all the fraud claims made are same irrespective of gender. Side collisions reported by females are fraudulent.

Similarly we do fill the '?' with mode of that column or by analysing the column. I had filled them with mode.

```
check('property_damage')
```

```
no. of unique values in property_damage is: 3  
unique values in property_damage is: ?      360  
NO      338  
YES     302  
Name: property_damage, dtype: int64
```



In Property damage '?' is replaced with YES. Since we observed that '?' have more fraud claims. By comparing with other data, claims with yes are more fraud. So we replaced that with YES.



```
dfc['property_damage']=dfc['property_damage'].replace('?', 'YES')

dfc['police_report_available'].value_counts()
NO      343
?        343
YES      314
Name: police_report_available, dtype: int64

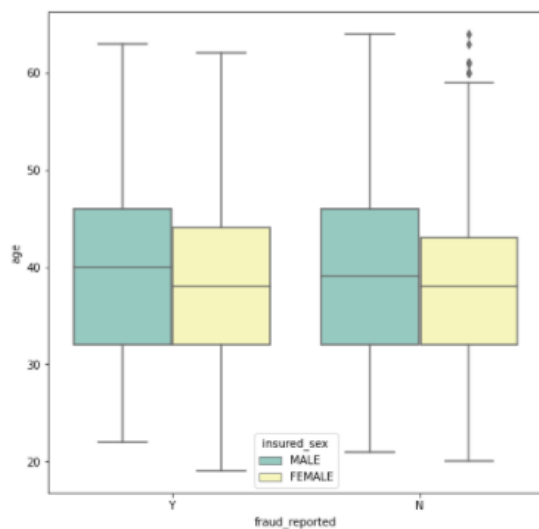
dfc['police_report_available']=dfc['police_report_available'].replace('?', 'NO')
```

Similarly police report is also filled with NO

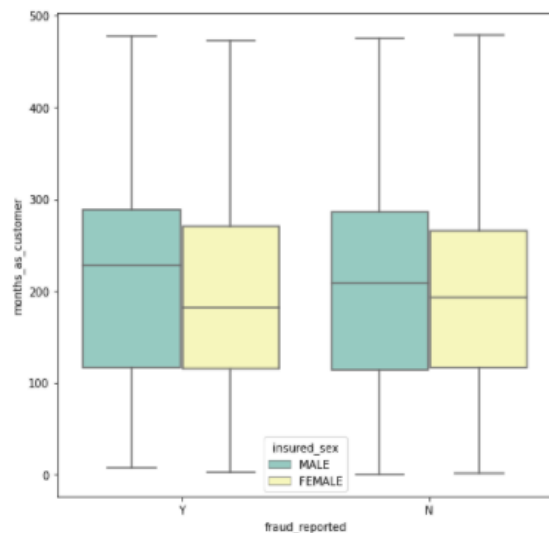
Let us see some important observations:

Average age and months being a customer: 38.9 yrs. old, 203.95 months as customer

```
Brief of age: count      1000.000000
mean       38.948000
std        9.140287
min        19.000000
25%        32.000000
50%        38.000000
75%        44.000000
max        64.000000
Name: age, dtype: float64
skewness: 0.47898804709224163
unique values count in age: 46
```

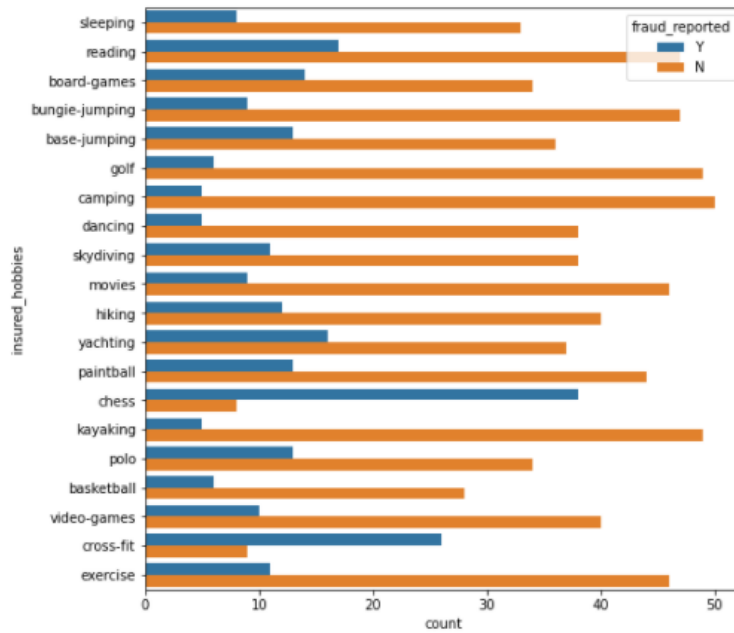


```
Brief of months_as_customer: count      1000.000000
mean      203.954000
std       115.113174
min        0.000000
25%       115.750000
50%       199.500000
75%       276.250000
max       479.000000
Name: months_as_customer, dtype: float64
skewness: 0.3621768477780205
unique values count in months_as_customer: 391
```

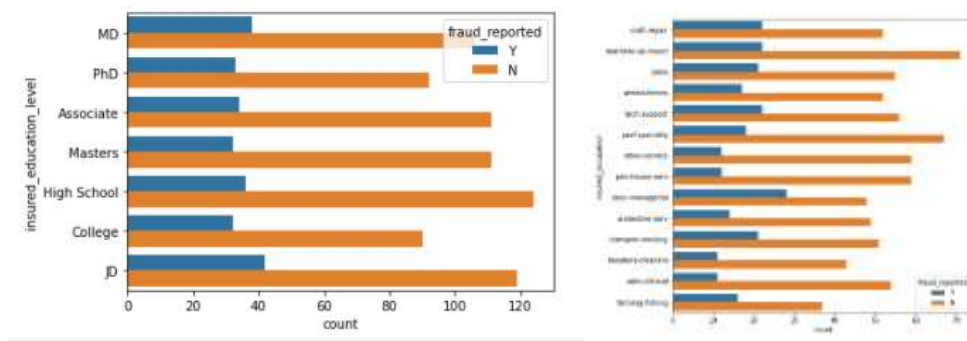


We can see males had done frauds from the age of 22/23 approx., females had done from teens, approx. from 18 yrs. Frauds committed by clients from their 1<sup>st</sup> month onwards.

Hobbies: People with hobby of chess and followed by cross-fit had done more fraud claims.

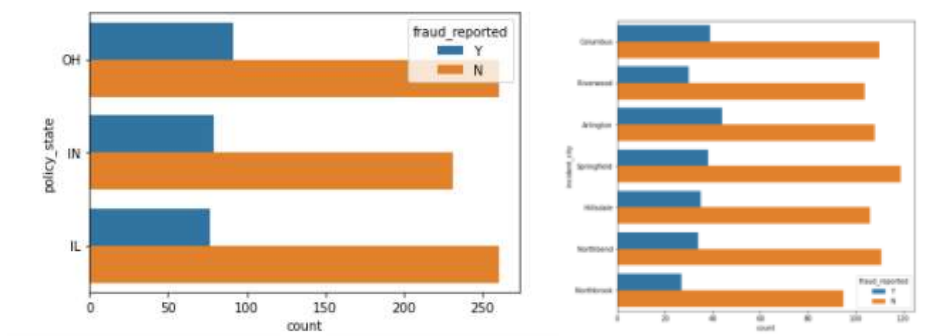


Education level: We can see, those who studied JD, MD has done more frauds comparatively.

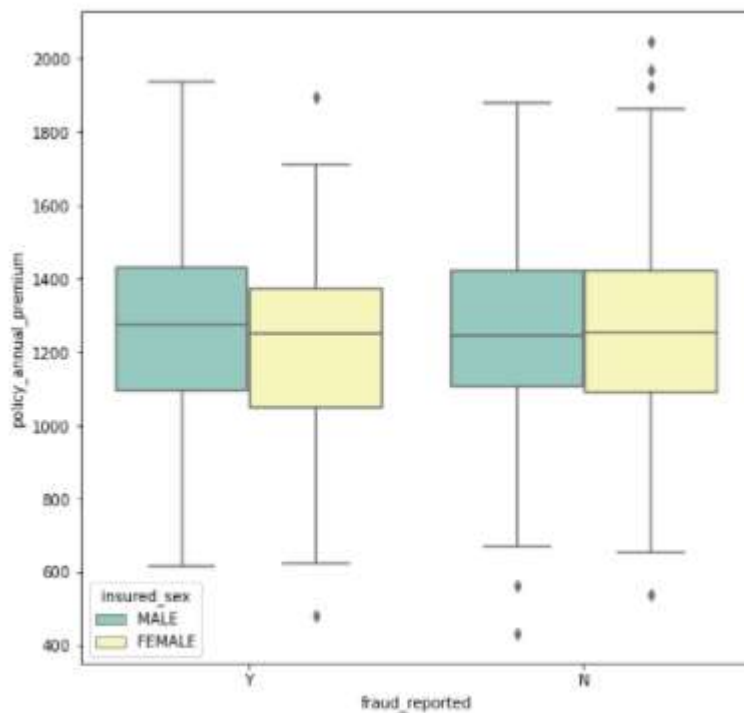


Occupation: People of job exec.manager are doing more frauds.

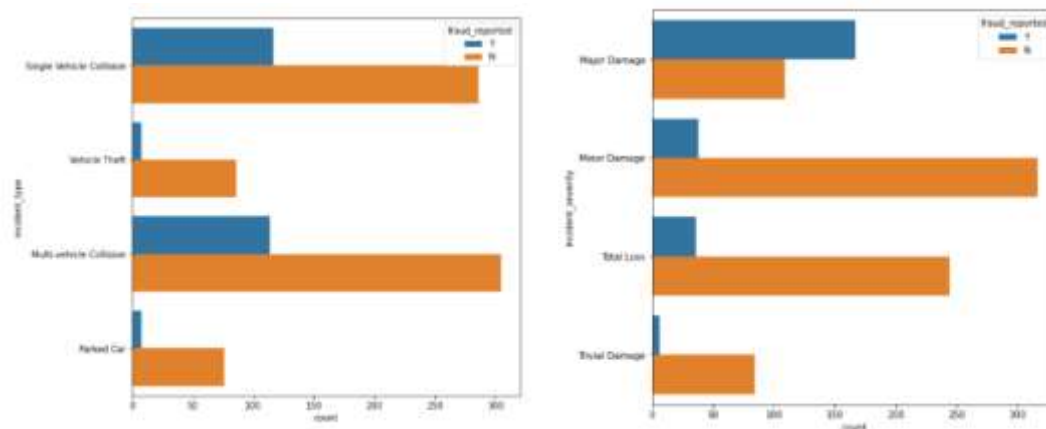
More frauds happened in state OH and in cities: Arlington, Columbus



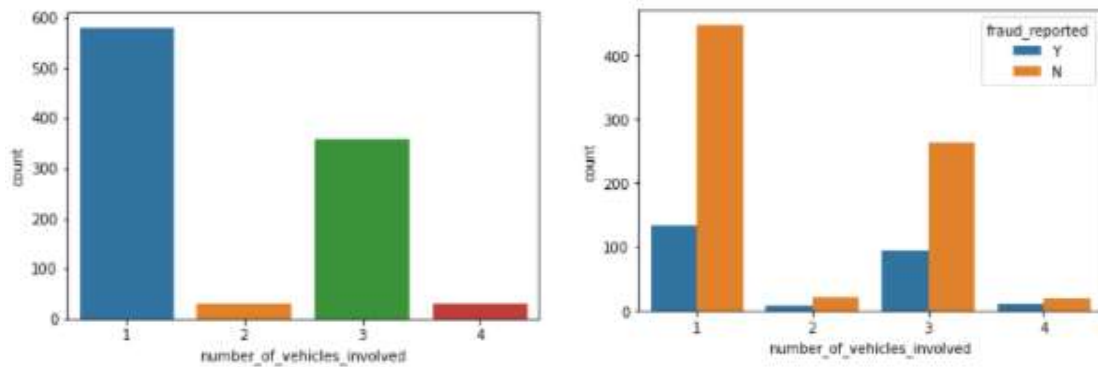
Policy premium: We seen that people with premium of above average are committing frauds



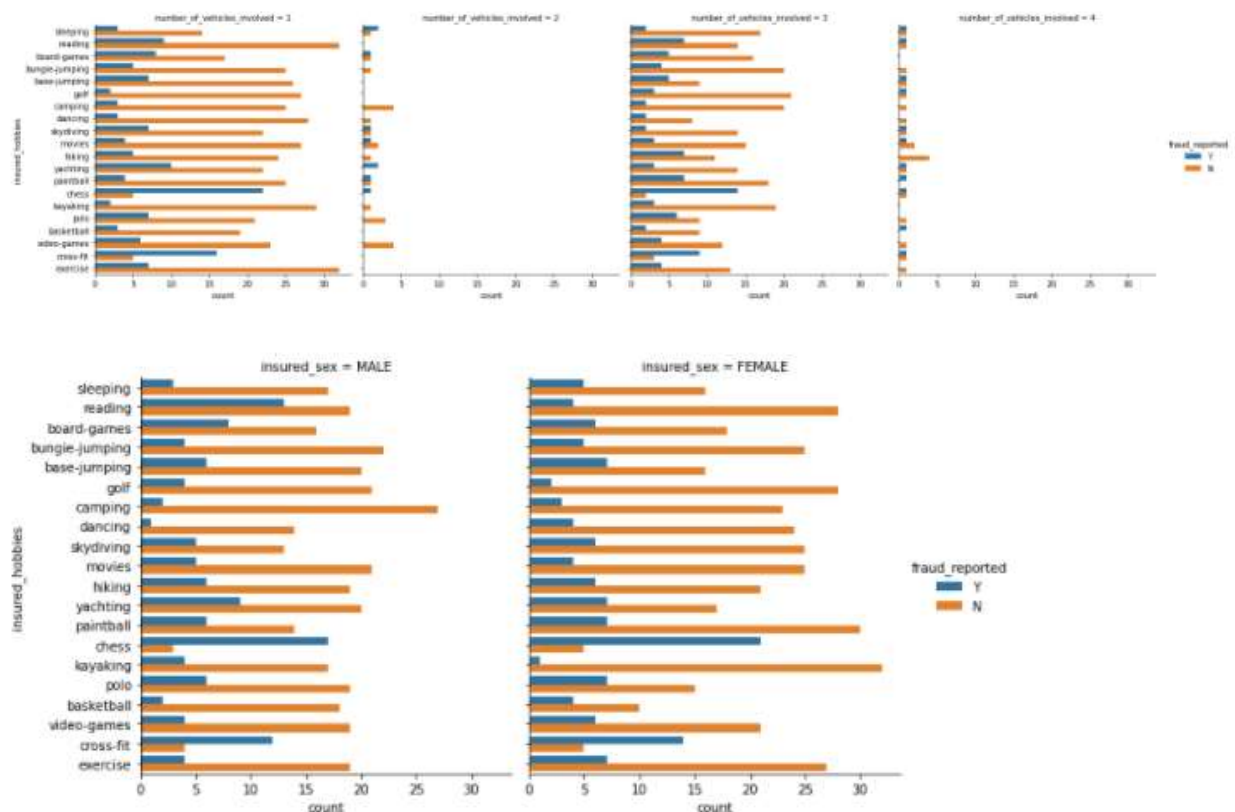
People who claimed incident type as Single vehicle collision or multiple vehicle and major damage collision are more fraudulent claims.

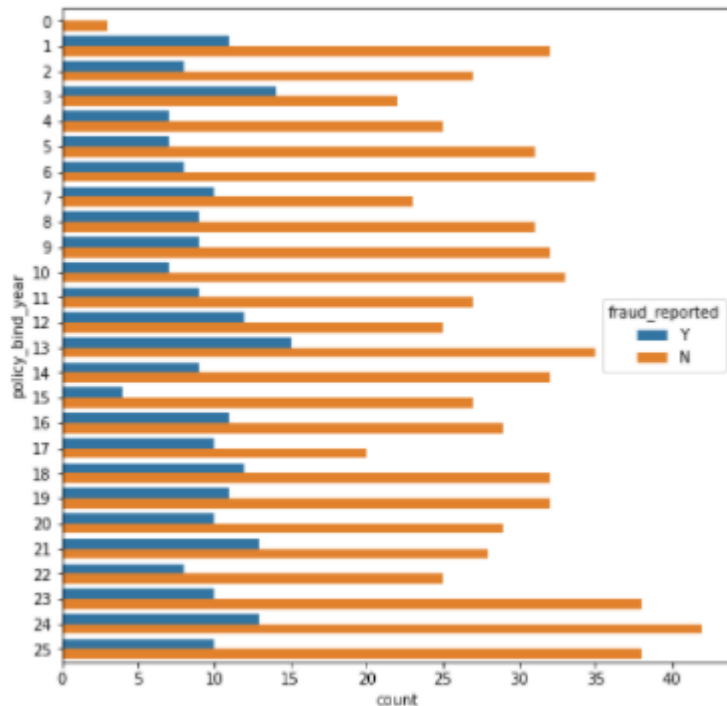


Majority of the incidents are 1 and 3 vehicles, also 1 involved are more fraud claims, but if we see carefully even though 2 and 4 vehicles involved claims are less, 50% of them are fraud.



Observing the fraud basis on no. of vehicles and Hobbies , Gender and Hobbies.





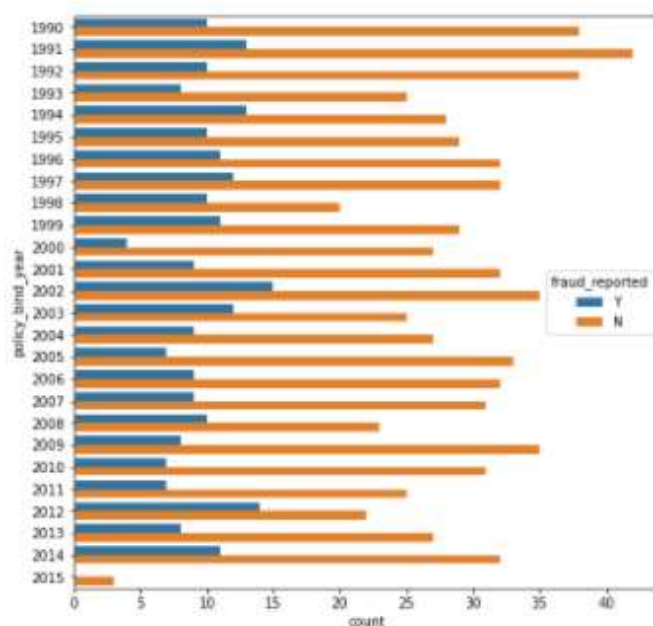
Based on number of years of policy binded

Similarly we observe and get insights from other columns.

We removed two more columns, incident date and incident location, which are not necessary, have almost unique values in each row.

We made a new column policy bind year from policy bind date and dropped the original column.

From that we observed policy been made from the years 1990 to 2015, and more frauds are of those who made insurance in year 1991, 2002, 2012.



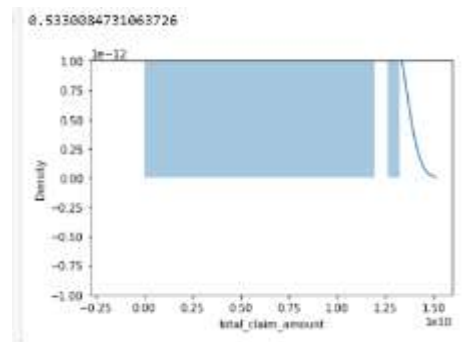
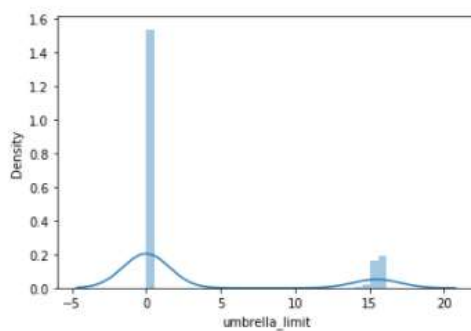
So to make it simple, we found the difference of number of years from 2015.

We made all the columns encoded and find the correlation.

Checking the skewness and reducing skewness in columns umbrella limit, total claim amount

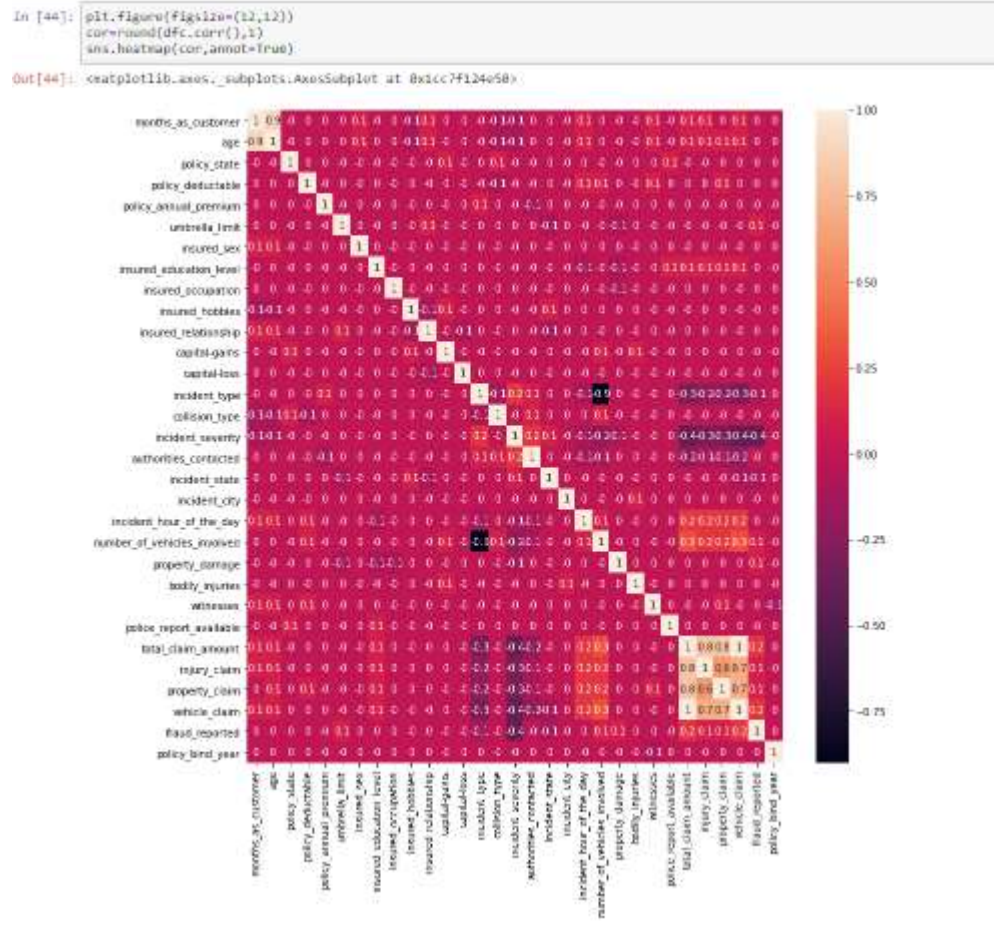
```
dfc.skew()

months_as_customer      0.362177
age                     0.478988
policy_state            -0.026177
policy_deductable       0.477887
policy_annual_premium   0.004402
umbrella_limit          1.806712
insured_sex             0.148630
insured_education_level -0.000148
insured_occupation      -0.058881
insured_hobbies          -0.061563
insured_relationship    0.077488
capital-gains           0.478850
capital-loss            -0.391472
incident_type           0.101507
collision_type          -0.033682
incident_severity       0.279016
authorities_contacted   -0.121744
incident_state          -0.148865
incident_city           0.049531
incident_hour_of_the_day -0.035584
number_of_vehicles_involved 0.502664
property_damage         -0.685977
bodily_injuries         0.014777
witnesses              0.019636
police_report_available 0.802728
total_claim_amount      -0.594582
injury_claim            0.264811
property_claim          0.378169
vehicle_claim           -0.621098
fraud_reported          1.175051
policy_bind_year        -0.052511
dtype: float64
```



Skewness reduced to 1.4 from 1.8, Skewness reduced to 0.5 from -0.59

## Plot of correlation:



## Observations:

- Age & months as customer are highly correlated, we will drop the column age.
- Total claim amount, injury claim, property claim, vehicle claim are correlated to each other. So we will drop all those columns except Total claim amount
- Incident hour of the day doesn't have any correlation with fraud. So we will drop it.
- There are many columns which have no correlation with fraud reported

Let us drop these columns, find the VIF to know the multicollinearity and drop other columns.

	variables	VIF factor
0	months_as_customer	0.101150
1	age	2.001211
2	policy_state	0.497810
3	policy_deductible	0.101017
4	policy_annual_premium	0.200311
5	umbrella_limit	1.000101
6	insured_sex	0.200101
7	insured_education_level	0.200101
8	insured_occupation	0.200101
9	insured_hobbies	0.200101
10	insured_relationship	0.200101
11	capital_gains	0.200101
12	capital_loss	0.200101
13	incident_type	0.200101
14	collision_type	0.200101
15	incident_severity	0.200101
16	authorities_contacted	0.200101
17	incident_status	0.200101
18	incident_city	0.200101
19	incident_hour_of_the_day	0.200101
20	number_of_vehicles_involved	0.200101
21	property_damage	0.200101
22	bodily_injuries	0.200101
23	witnesses	0.200101
24	policy_report_available	0.200101
25	total_claim_amount	0.200101
26	injury_claim	0.200101
27	property_claim	0.200101
28	vehicle_claim	0.200101
29	fraud_reported	0.200101
30	policy_lend_year	0.200101

By viewing VIF I dropped columns

Number of vehicles involved, premium deductables, capital gain and capital loss.



Finally two last steps before making model, Split x & y, then scale the x and finally balance the dataset.

```
In [56]: #Let us split the columns x & y
x=dfc.drop('fraud_reported',axis=1)
y=dfc['fraud_reported']

In [58]: #Scaling the dataset
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_scaled=sc.fit_transform(x)

In [59]: from imblearn.over_sampling import SMOTE
smt=SMOTE()
x_balanced,y_balanced=smt.fit_resample(x_scaled,y)
```

X is features, y is our target.

Found best random state & split.

```
#Let us import and split x y at best random state
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix,accuracy_score,classification_report
from sklearn.model_selection import train_test_split

lo=LogisticRegression()
rs=0
acsc=0

for i in range(1000):
    x_train,x_test,y_train,y_test=train_test_split(x_balanced,y_balanced,test_size=0.2,random_state=i)
    lo.fit(x_train,y_train)
    pred=lo.predict(x_test)
    acc=accuracy_score(y_test,pred)
    if acc>acsc:
        acsc=acc
        rs=i
print(f'Best score:{acsc}\n random state: {rs}')
```

Best score:0.813953488372093  
random state: 832

```
x_train,x_test,y_train,y_test=train_test_split(x_balanced,y_balanced,test_size=0.2,random_state=832)
```

Imported different models to try and select the best one

```
#Importing other models:
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
models=[LogisticRegression(),SVC(),DecisionTreeClassifier(),RandomForestClassifier(),GradientBoostingClassifier(),KNeighborsClassifier()]
for m in models:
    m.fit(x_train,y_train)
    predm=m.predict(x_test)
    print(f'{m}:')
    print('accuracy score:',accuracy_score(y_test,predm))
    print('confusion matrix:\n',confusion_matrix(y_test,predm))
    print('classification report:\n',classification_report(y_test,predm))
    cvscore=cross_val_score(m,x_balanced,y_balanced,cv=5)
    print('mean cv score:',cvscore.mean())
    print('\n')
```

From above we got results of accuracy score, f1 score, evaluation report and cross validation score of different models –

Model	F1 score	Cross validation	Difference
Logistic Regression	81	74.5	6.5
Decision Tree Classifier	85	84.9	0.1
Random Forest Classifier	91	87	4
Gradient Boosting Classifier	91	88	3
K Neighbors Classifier	72	73	1
SVC	89	84.5	4.5

From above we can see, Decision Tree is the best model with least difference between CV score and F1 score. These are the metrics generated:

```
DecisionTreeClassifier():
accuracy score: 0.8538205980066446
confusion matrix:
[[129  21]
 [ 23 128]]
classification report:
              precision    recall  f1-score   support

     0       0.85         0.86         0.85         150
     1       0.86         0.85         0.85         151

 accuracy          0.85         0.85         0.85         301
 macro avg         0.85         0.85         0.85         301
 weighted avg      0.85         0.85         0.85         301

mean cv score: 0.8491052048726466
```

From confusion matrix we can see

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP <b>21</b>
	Negative (0)	FN <b>23</b>	TN

Our model is saying 23 members as not fraud where they are actually frauds. We need to reduce this for better results.

FP is ok because checking whether fraud or not is less cost to company than paying huge amount to the original fraud claim.

In our case we need least False Negative.

But wait there is almost 6% difference in Decision Tree classifier, Random Forest or Gradient Boosting. Let us observe their metrics once

If we see other models: It has 3% difference but giving the least FN

```
GradientBoostingClassifier():
accuracy score: 0.9136212624584718
confusion matrix:
[[133 17]
 [ 9 142]]
```

```
classification report:
              precision    recall  f1-score   support

     0       0.94      0.89      0.91       150
     1       0.89      0.94      0.92       151

 accuracy          0.91
 macro avg         0.91
 weighted avg      0.91
```

```
mean cv score: 0.8803521594684385
```

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP <b>17</b>
	Negative (0)	FN <b>9</b>	TN

By tuning the decision tree classifier: The model accuracy increased from 85 to 87% only. With FN as 21.

```
: from sklearn.model_selection import GridSearchCV
diff_para={'criterion':['gini', 'entropy'],'splitter':['best', 'random'],'random_state':[10,20,30,40,50,60,70]}
grid=GridSearchCV(estimator=dtc,param_grid=diff_para)
grid.fit(x_train,y_train)
print('Best Score:',grid.best_score_)
print('Best Estimator:',grid.best_estimator_)
print('Best Parameter:',grid.best_params_)
```

```
Best Score: 0.8395539419087136
Best Estimator: DecisionTreeClassifier(criterion='entropy', random_state=60)
Best Parameter: {'criterion': 'entropy', 'random_state': 60, 'splitter': 'best'}
```

```
: dtc=DecisionTreeClassifier(criterion='entropy', random_state=60,splitter='best')
dtc.fit(x_train,y_train)
pred=dtc.predict(x_test)
print('Metrics of GradientBoostingClassifier:')
print('Score:',accuracy_score(y_test,pred))
print('Confusion Matrix:',confusion_matrix(y_test,pred))
print('Classification Report:',classification_report(y_test,pred))
```

```
Metrics of GradientBoostingClassifier:
Score: 0.867109634551495
Confusion Matrix: [[131 19]
 [ 21 130]]
Classification Report:
              precision    recall  f1-score   support

     0       0.86      0.87      0.87       150
     1       0.87      0.86      0.87       151

 accuracy          0.87
 macro avg         0.87
 weighted avg      0.87
```

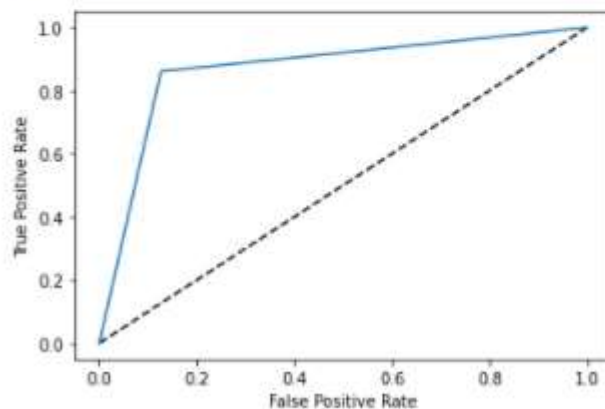
## AUC-ROC:

```
3]: #Auc Roc curve

from sklearn.metrics import roc_auc_score
#Predicting the probability of having 0 in the x-test
y_pred_prob=dtc.predict_proba(x_test)[:,-1]
y_pred_prob

#Visualising
from sklearn.metrics import roc_curve
fpr,tpr,thresholds=roc_curve(y_test,y_pred_prob)

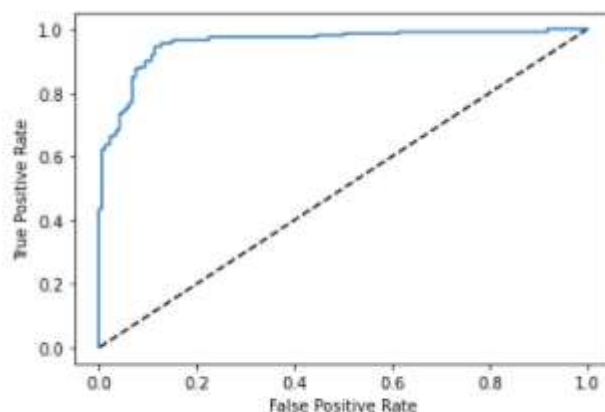
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr,label='KNeighborsClassifier')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
auc_score=roc_auc_score(y_test,dtc.predict(x_test))
print('Score:',auc_score)
```



Score: 0.867130242825607

For the Gradient Boosting model, even after tuning there is no improvement in accuracy:

## AUC-ROC:



Score: 0.9135320088300221

I am considering this as my best model and saving it to predict.

Saving the model:

```
import joblib
joblib.dump(gc, 'Auto_Insurance_Fraud.obj')

['Auto_Insurance_Fraud.obj']
```

Conclusion:

Now we made a model, that can predict whether the claim made is Fraud or not with a 91% accuracy.

We have Seen, Gender, number of months as customer, Premium, Hobbies, Education , Type of incident , Incident severity and the policy bind year are the most important columns that help us to predict and understand whether the claim is fraud or not.

- Customers with education of JD are doing more fraud.
- Customers with hobbies of CHESS, CROSS-FIT are doing more fraud.
- Customers who claimed the incident to be single vehicle collision or Multi vehicle collision are more fraud claims.
- Gender is almost same in both the cases of fraud claims
- Premium – Customers who opted for higher premiums done frauds
- Majority of claims made by customers with policy bin year 1991, 2002, and 2012 are frauds.

We encoded the values, categorical columns. Our result is also encoded as 0 and 1 where

0 is NO

1 is YES

Later it will be decoded, when the model is deployed.

Attaching the GitHub link along with this, where you can get dataset, the complete step by step procedure followed jupyternotebook .ipynb file.

[https://github.com/Makham-Sai-Girish/Auto\\_Insurance\\_Fraud-Claim\\_detection](https://github.com/Makham-Sai-Girish/Auto_Insurance_Fraud-Claim_detection)